Geolocation Prediction from Jodel and Twitter Messages

Karl Oskar Magnus Holm
Engineering Science and ICT - Department of Geomatics / 2023
`koholm@stud.ntnu.no`

Abstract

This paper takes on the shared tasks on social media variety geolocation from the VarDial workshops in 2020 and 2021, focusing on the subtask on geolocation of Swiss Jodel messages. Both year's winners used BERT Transformer models, and this paper builds upon their work, investigating if newer language-specific models, other map projections, or different hyperparameters can improve the accuracy. While I was unable to match the best results from 2020 and 2021, it is clear from my results that relative improvement from the BERT models used in 2020 and 2021 is outperformed by newer language-specific models. Language-specific variants of Google's BERT and Meta's X-Mod both showed significant improvements over older the models on the task of social media variety geolocation.

1 Introduction

This project is based on the shared task on Social Media Variety Geolocation (SMG) from VarDial 2020 and 2021 (seventh and eighth edition, respectively), the Workshop on Natural Language Processing (NLP) for Similar Languages, Varieties and Dialects (Chakravarthi et al., 2021; Gaman et al., 2020). The aim of the task differs somewhat from the most common types of NLP VarDial tasks, where the goal typically is to choose from a finite set of variety labels (Scherrer and Ljubešić, 2021, p. 1). Here, the goal is to predict a set of scalars, namely the latitude and longitude from which a social media post was posted. This VarDial task stayed the same from 2020 to 2021, including three language areas: the Bosnian-Croatian-Montenegrin-Serbian language area, the German language area (Germany and Austria in this case), and the German-speaking Switzerland.

This project is limited to the latter of these language areas, that is, the German-speaking Switzerland. Reasons for this include the limited time scope of the task, and having to share the necessary computing resources with fellow students at the department. The goal is to try and recreate the results of Scherrer and Ljubešić (2020) who used a BERT-based classifier, making it a double regression task. I focus on the 2020 dataset because there were a lot more submissions this year as opposed to in 2021, due to the short time between the announcement of the shared task and the submission deadline (Chakravarthi et al., 2021, p. 6).

The reason for picking the task on the German-speaking Switzerland is its similarities to the dialectal landscape of Norway. Røyneland (2009, p. 14) writes that "dialects in Norway have had and still have a much stronger position than dialects ... in most of Europe, except for the German-speaking part of Switzerland.". I therefore find it reasonable to assume that a method which works well on the Swiss dataset could also perform well on a Norwegian dataset. Unfortunately, I was unable to find a suitable Norwegian dataset, and creating one proved too difficult and time-consuming to be worthwhile, considering the time horizon of the project.

Code and data used to develop models is available at https://github.com/oskarhlm/TDT13.

2 Background

This section will elaborate on technologies central to this project. It is assumed that the reader has basic understanding of what Language Models (LMs) are, and also that they are somewhat wandered in the world of Natural Language Processing (NLP).

2.1 The Transformer Architecture

Vaswani et al. (2017) managed to achieve new state-of-the-art results for machine translation tasks with their introduction of the Transformer architecture. The Transformer has later been proved effective for numerous downstream tasks, and for a variety of modalities. Titleing their paper 'Attention Is All You Need', Vaswani et al. suggest that their attention-based architecture renders Recurrent Neural Networks (RNNs) redundant, due to its superior parallellization abilities and the shorter path between combinations of position input and output sequences, making it easier to learn long-range dependencies (Vaswani et al., 2017, p. 6).

The Transformer employs self-attention, which enables the model to draw connections between arbitrary parts of a given sequence, by-passing the long-range dependency issue commonly found with RNNs. An attention function maps a query and a set of key-value pairs to an output, calculating the compatibility between a query corresponding key (Vaswani et al., 2017, p. 3). Looking at Vaswani et al.'s proposed attention function (1), we observe that we take the dot product between the query Q and the keys K, where Q is the token that we want to compare all the keys to. Keys similar to Q will get a higher score, e.g. be *more attended to*. These differences in attention is further emphasized be applying the softmax function. The final matrix multiplication with the values V, being the intitial embeddings of all input tokens, will give us a new embedding in which all tokens have some context from all other words. We improve the attention mechanism by mutiplying each query, key, and value with weigth matrices learned through backpropagation. Self-attention is a special kind of attention in which queries, keys, and values are all the same input sequence.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \tag{1}$$

Attention blocks can be found in three places (Vaswani et al., 2017, p. 5) in the Transformer architecture (I will use machine translation as example, say, from Norwegian to German):

1. In the encoder block to perform self-attention on the input sequence (which is in Norwegian)

2. In the decoder block to perform self-attention on the output sequence (which is in German)

3. In the decoder block to perform cross-attention (or encoder-decoder attention) where each position in the decoder attends to all positions in the encoder

The Transformer established a new state of the art in machine translation Vaswani et al. (2017), and is the fundamental building block of LMs like BERT.

2.2 BERT

Bidirectional Encoder Representations from Transformers (BERT) is a family of language models which was first introduced in 2018 and is designed to facilitate a wide range of downstream tasks (Devlin et al., 2019, May, p. 5). The BERT architecture consists of stacked bidirectional Transformer encoders. The self-attention mechanism coupled with a masked language modelling pre-training step allows for training of deep bidirectional representations. 15 percent of words are masked with the special `[MASK]` token during this pre-training step and left for the model to predict (Devlin et al., 2019, May, p. 4). The second of the two unsupervised tasks used during pre-training, is Next Sentence Prediction (NSP), where the special `[CLS]` (found at the start of each tokenized sequence) is used to predict if a senctece `B` follows `A`. The input representation then looks like this:

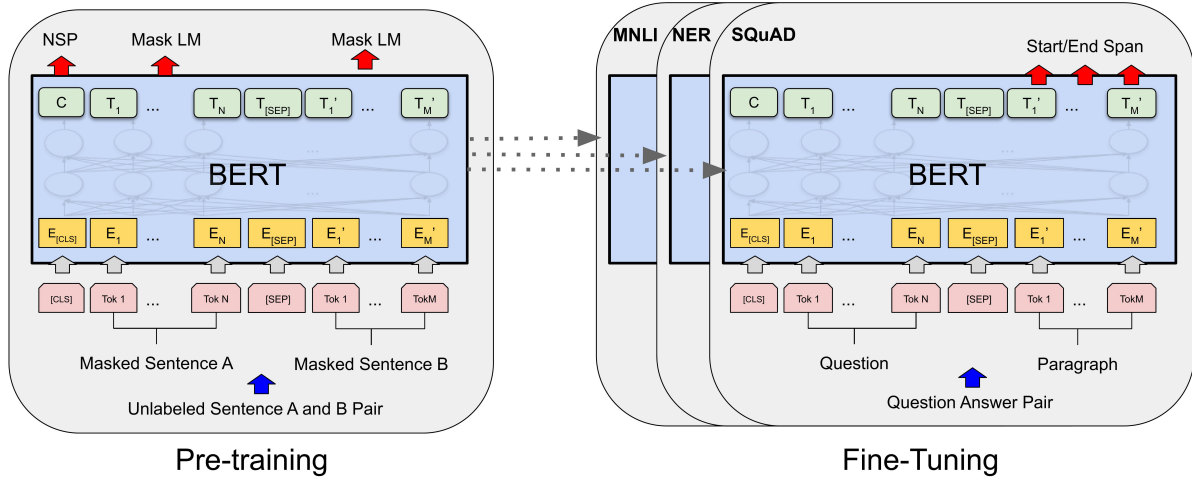`[CLS]` this is sentence A `[SEP]` and this is sentence B `[SEP]`

Figure 1: High-level overview of the pre-training and fine-tuning procedures for BERT (Devlin et al., 2019, May, p. 3)

Figure 1 shows a high-level overview of the pre-traning and fine-tuning procedures. For the pre-training, each token of the input sequence, consisting of a sentence pair and the classification token, [CLS], is transformed into embeddings (vector representations). These per-token embeddings include information of the meaning of the word itself, the meaning of the sentence/segment it belongs to, and the token's position in the full input. These embeddings are the pass through a stack of Transformer encoders (12 and 24 for **BERT**$_{\text{BASE}}$ and **BERT**$_{\text{LARGE}}$, respectively), allowing the model to learn more complex patterns and of different granularities (token, sentence, document).

Fine-tuning

2.3 X-Mod

Cross-lingual Modular (X-Mod) models (Pfeiffer et al., 2022, July) attempt to tackle the common problem of multi-linguality in language models. Typically, when one attepmts to train a language model be multilingual by training on numerous languages, the performance tends to drop after reaching a certain level of performance - *the curse of multilinguality* (Pfeiffer et al., 2022, July, p. 1).

2.4 Geodesic Terminology and Metrics

The evaluation is based upon the Haversine formula, with the Earth's radius is assumed to be 6371 km. The evaluation metric is the median Haversine distance (2) between the predicted coordinates and the ground thruth (Scherrer and Ljubešić, 2020, p. 4). A formulation of the Haversine distance can be found on its Wikipedia page[1] where it is described as "the great-circle distance between two points on a sphere given their longitudes and latitudes". The distance d can be expressed as

$$d = 2r \arcsin \left(\sqrt{\sin^2 \left(\frac{\phi_2 - \phi_1}{2} \right) + \cos(\phi_1) \cos(\phi_2) \sin^2 \left(\frac{\lambda_2 - \lambda_1}{2} \right)} \right) \tag{2}$$

where ϕ and λ are latitude and longitude values.

Different map projections were used in the project. The Universal Transverse Mercator (UTM) map projection splits Earth's surface into 60 zones in the latitudal direct and 19 zones in the longitudal direction, forming a grid. Doing this allows us to express coordinates in meters within a grid zone and still obtaining high accuracy measurements. Coordinate values in UTM lie in the six figures, with the easting of the central meridian[2] being defined to 500 000 meters to avoid negative easting values within the zone.

The Swiss coordinate system, LV95 (Federal Office of Topography swisstopo, n.d.), was also explored. Its center coordinates are defined to the Swiss capital of Bern and the values lie in the seven figures.

---

[1] https://en.wikipedia.org/wiki/Haversine_formula
[2] https://gisgeography.com/central-meridian/

	lat	lon	text
0	47.22	7.43	Dr Chester Bennington isch tot (pensive face)(pensive face)(pensive face) #rip #linkinpark Dr Manager heds bestätiged (expressionless)...
1	46.86	8.21	Mini Fründin hed Lust uf Doktorspieli gha... ... sie hocked jetzt sit 2 Stund...
2	47.39	8.18	Slayer isch besser. Det han ich gescht mini Drohne stiege lah (smiling face with smiling eyes) Cool was hesch f...
3	47.37	8.78	gaht au innere stund? bin grad am speck brate (nerd face) So langt chunsch ja münd eifach...
4	47.39	8.04	sie: thy er: ? sie: thy= thank you er: player sege thx...

Table 1: The first five rows of the training dataset

3 Related Work

This work builds on top of the work of Scherrer and Ljubešić (2020) and Scherrer and Ljubešić (2021). They were the only participants in VarDial who used a large LMs like BERT in the shared task on social media variety geolocation, and did so with great success, winning the shared task in both 2020 and 2021. Scherrer and Ljubešić converted the task into a double regression problem, where latitude and longitude values are predicted from the output of a large LM. They experimented with different pretrained models, coordinate encodings, and hyperparameters. Their main finding was that single-language models outperform multilingual models, the latter of which perform worse due to capacity dilution and tokeniers yielding suboptimal text splitting (Scherrer and Ljubešić, 2020, p. 3). As they were unable to find pre-trained model a pre-trained model for Swiss German they instead trained `bert-base-german-uncased`[3] (German BERT) on the SwissCrawl corpus (Linder et al., 2020, June). Training a total of 48 models, Scherrer and Ljubešić were able to achieve a median distance of 15.72 km in this unconstrained setting using the default data split. They got a median distance of 15.45 km by using a substantioal portion of the development set for training (Scherrer and Ljubešić, 2020, p. 6).

Gaman et al. (2020) and Chakravarthi et al. (2021) summarize the findings in the 2020 and 2021 editions of VarDial, including attempts made on the Social Media Variety Geolocation (SMG) task. While Scherrer and Ljubešić (2020) generally dominated the leaderboards, Benites de Azevedo e Souza et al. (2020) proposed a method that performed best among constrained submissions on the Swiss task (Gaman et al., 2020, pp. 8–9), and only marginally worse than Scherrer and Ljubešić's unconstrained submissions. Benites de Azevedo e Souza et al. (2020) use K-Means clustering (Lloyd, 1982) of locations and predicting cluster identities, framing the problem as a classification task rather than a regression task. Their best submission extracts features from different levels of token granularity, training a separate SVM for each feature set, before feeding the distances to the decision boundaries for each feature classifier as input to a SVM meta-classifier.

4 Datasets

Data from VarDial 2020 and 2021 was acquired from a GitHub repository[4] created by Yves Scherrer's, co-author of the winning solution for the SMG task, both years. All but the test dataset has a ground truth associated with it, and I assume this unlabelled test dataset was used for a private leaderboard. The training, development, and test gold datasets have 22600, 3097, and 3068 labelled samples, respectively. Table 1 shows the first five rows of the training dataset. It was collected by Hovy and Purschke (2018, October, pp. 2–3) using the (at the time) publicly available Jodel API.

While Switzerland is a country with four official languages (Swiss-German, French, Italian, and Ro-

---
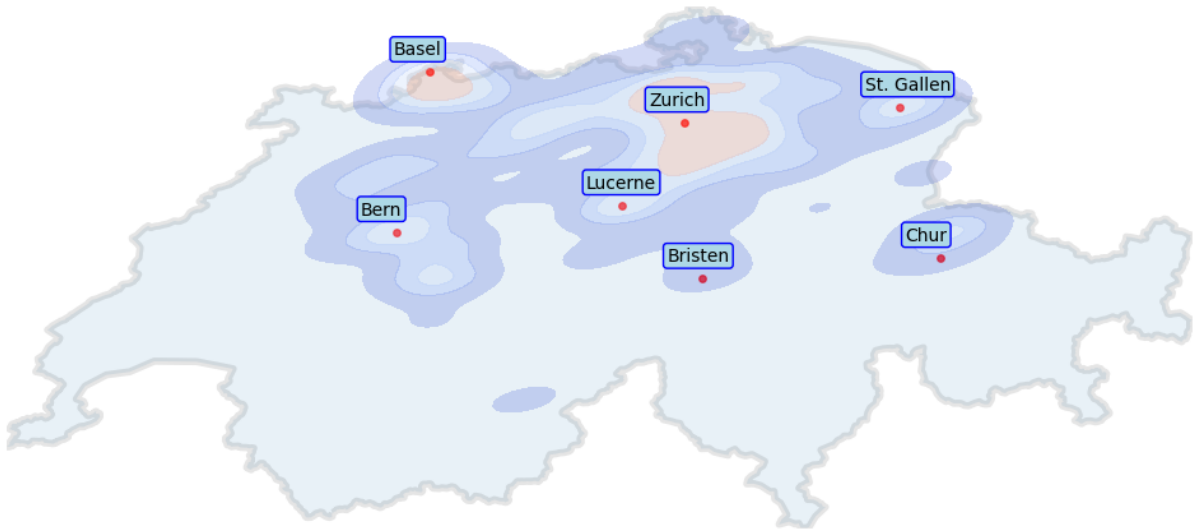
[3] https://huggingface.co/dbmdz/bert-base-german-uncased
[4] https://github.com/yvesscherrer/vardial-shared-tasks

Figure 2: Heatmap of the training data

Model Name	Model Type
dbmdz/**bert-base-german-uncased**	BERT
statworx/**bert-base-german-cased-finetuned-swiss**	BERT
ZurichNLP/**swissbert**	X-Mod

Table 2: Pre-trained models used in the project

mansh Grishum), the dataset contains only Swiss-German Jodel messages, focusing on dialectal differences (the "Dial" in VarDial). This results in the heatmap in Figure 2.

Attepmts were made to aquire a Norwegian dataset based on Norwegian Twitter/X messages using the method described in Ljubešić et al. (2016) but due to recent changes in the Twitter/X API[5] I was unable to do this. I also explored the possibility of using Norwegian Jodel messages but to my knowledge their API is no longer available to the public.

5 Model

Figure 3 shows a rough model architecture. It consists of a pre-trained BERT model with a classification head on top. This architecture is only representative for the BERT-based models (all but one).

The pre-trained models that were tested in this project are listed in Table 2. When selecting models I was mostly interested in those that are trained on Swiss corpera, seeing as this proved important in Scherrer and Ljubešić (2020). They were utilized through Huggingface's `transformers` interface.

The classification head has two outputs: the latitude coordinate and the longitude coordinate. It takes as input the output corresponding to the `[CLS]` token, which captures the aggregated sequence representation. The hyperbolic tangent activation function adds some non-linearity to the output before it fed into a fully connected linear layer, where latitude and longitude values are predicted.

`bert-base-german-uncased` is the same model that Scherrer and Ljubešić (2020) used for their best solutions. It is trained of a Wikipedia dump, EU Bookshop corpus, and more. `bert-base-german-cased-finetuned-swiss` is based upon `bert-base-german-cased` and is finetuned on the Leipzig Corpora Collection and SwissCrawl. `ZurichNLP/swissbert` is the only non-BERT model used. It is rather based on X-Mod, and has adapters trained for German, French, Italian, and Ramansh Grishun.

_____
[5]https://www.theverge.com/2023/5/31/23739084/twitter-elon-musk-api-policy-chilling-academic-research
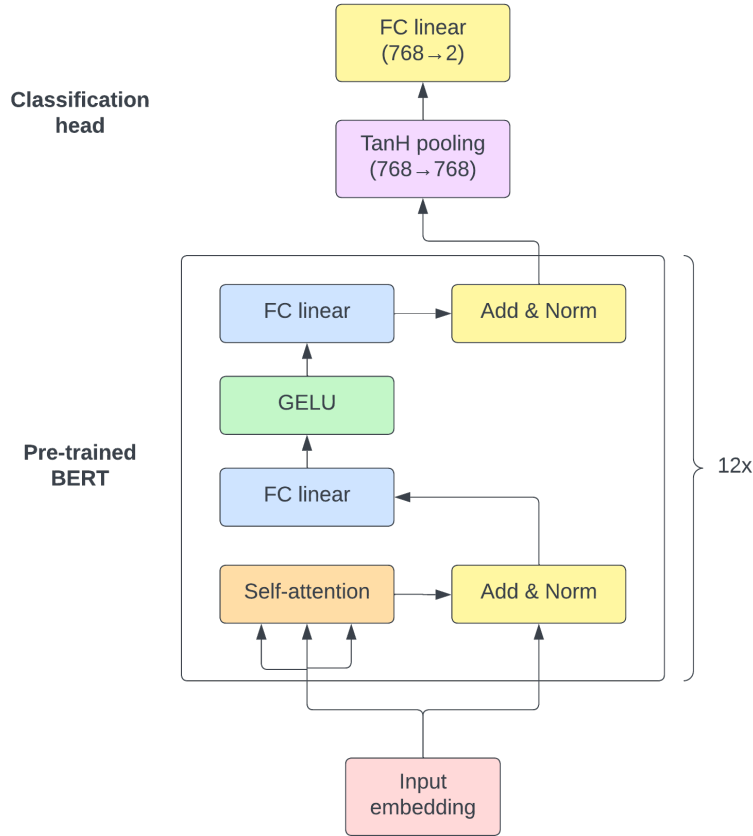
Figure 3: Model architecture

6 Experiments and Results

In this section I will elaborate on my approach when training models, before presenting the most important findings from my experiments.

6.1 Experimental Setup

All experiments were performed on an NVIDIA GeForce RTX 4090, having 24 GB G6X memory[6]. Computing resources belong to the Department of Geomatics at NTNU and are shared with fellow 5th year geomatics students. PyTorch[7] was used to create a training loop, and Huggingface's `trans-formers` library was used to fetch pre-trained models from their hub.

The best results of Scherrer and Ljubešić (2020) came from using a language-specific BERT. As no pre-trained model was found in 2020, they fine-tuned the pre-trained `bert-base-german-uncased`[8] model on the SwissCrawl corpus (Scherrer and Ljubešić, 2020, pp. 3–4). Since then, a pre-trained Swiss BERT model[9] has been released, which I used in my experiments.

Becuase of the limited timespan and computing resources of this project, I opted to freeze certain of Scherrer and Ljubešić's hyperparameters. This includes the maximum sequence length and the batch size, the latter of which was also limited by the GPU memory. The loss function (MAE/L1) and scaler (joint (Scherrer and Ljubešić, 2020, p. 5)) also largely remained unchanged. The focus of my experiments was to compare performance of pre-trained model, and seeing what effect different coordinate projections and learning rates have.

---

[6]https://www.nvidia.com/nb-no/geforce/graphics-cards/40-series/rtx-4090/
[7]https://pytorch.org/
[8]https://huggingface.co/bert-base-german-uncased
[9]https://huggingface.co/statworx/bert-base-german-cased-finetuned-swiss

Pre-trained model	Coordinate Projection	LR/Scheduler	Median Distance [km]
dbmdz/**bert-base-german-uncased**	lat/lon	4e-5	17.81
statworx/**bert-base-german-cased-finetuned-swiss**	lat/lon	4e-5	17.08
statworx/**bert-base-german-cased-finetuned-swiss**	UTM	ReduceLROnPlateau	16.52*
statworx/**bert-base-german-cased-finetuned-swiss**	UTM	2e-5	16.05*
statworx/**bert-base-german-cased-finetuned-swiss**	lat/lon	2e-5	16.19*
statworx/**bert-base-german-cased-finetuned-swiss**	LV95	2e-5	**15.76***
ZurichNLP/**swissbert**	UTM	2e-5	17.59*

Table 3: Highlighted results

* Proportion of developmentset used as additional samples for training

6.2 Experimental Results

A total of 15 models were trained. Table 3 shows a selection of the most interresting results along with to coordinate projection and learning rate/learning rate scheduler used. A full overview of configurations and their corresponding results can be found on the project's GitHub page. Fine-tuned models and training logs can be found in this Google Drive folder.

The best results were achieved when using the `statworx/bert-base-german-cased-finetuned-swiss` pre-trained model with a reduced development/validation dataset. Learning rate schedulers did not prove very efficient for this task and a low learning rate of 2e-5 yielded the best results instead. No real differnece in performance was observed when using the UTM projection over raw latitude/longitude values, but the Swiss-native LV95 projection gave the model a performance boost into the 15-kilometer range. The X-Mod based `swissbert` model did not prove efficient for the task at hand.

Figure 4 show the pointwise distance from the ground truth when using the best-performing model to make predictions on the test gold dataset.

Include the correct figure

7 Discussion

Results show that the Swiss LV95 map projection proved most efficient for predicting geolocation from Swiss Jodel messages. This may indicate that using a metric Coordinate Reference System (CRS) like LV95 over a spherical representation like latitude and longitude value can be beneficial in a double regression task on predicting geographical coordinates. It could seem that the model finds it harder to learn spherical representations. Using the metric UTM projection also yielded minor improvements over latitude/longitude coordinate representations. These findings are counter to those of Scherrer and Ljubešić (2020, p. 5), who found that raw latitude and longitude does not perform worse.

Not surprisingly, the language-specific model (`bert-base-german-cased-finetuned-swiss`) proved more suitable this task. Being trained on the Leipzig Corpora Collection (Goldhahn et al., n.d.) and SwissCrawl (Linder et al., 2020, June), its creators were able to improve show a 5 percent improvement over its German parent model (`bert-base-german-cased`[10]). It seems this fine-tuning enhances the model's ability to pick up on dialectal details in the data. The X-Mod-based `swissbert` model, which is based on a model designed to be multilingual (Pfeiffer et al., 2022, July),

[10]https://huggingface.co/bert-base-german-cased

Figure 4: Pointwise distance from the ground truth

did not seem to possess the same dialectal knowledge, and performed only marginally better than the German `ber-base-german-uncased` model.

Furthermore, it is clear from the results that the learning rate schedulers used did not improve the test score. `ReduceLROnPlateau` and `OneCycleLR` were tested and while they were able to greatly reduce the convergence time, they were unable to achieve satisfiable median distances. Since the schedulers showed such little promise for this task they were not investigated much further. I do think, however, that they could prove efficient if one can find a suitable set of initialization parameters.

8   Conclusion and Future Work

This paper addressed the shared task on Social Media Variety Geolocation from VarDial 2020 and 2021, and shows that newer language-specific Language Models are very capable of predicting latitude/longitude pairs from text samples, obtaining a median error or 15.76 kilometers on the Swiss task. It also suggests that using metric Coordinate Reference Systems could be better for what becomes a double regression problem in the geographic realm. The best model achieved what would have been a second place in the original competition.

Further reasearch should look into using classifier based upon the **BERT$_{\text{LARGE}}$** base model. As of 22nd November 2023there is no such model fine-tuned on Swiss corpera known to the author of this paper. Being twice as large as the **BERT$_{\text{BASE}}$** model it should be able to pick up on finer dialectal details, if fine-tuned.

References

Benites de Azevedo e Souza, F., Hürlimann, M., von Däniken, P., & Cieliebak, M. (2020). ZHAW-InIT : Social media geolocation at VarDial 2020. *Workshop on NLP for Similar Languages, Varieties and Dialects, Barcelona (Spain), Online, 13 December 2020*, 254–264
Accepted: 2021-02-04T13:13:06Z.

Chakravarthi, B. R., Mihaela, G., Ionescu, R. T., Jauhiainen, H., Jauhiainen, T., Lindén, K., Ljubešić, N., Partanen, N., Priyadharshini, R., Purschke, C., Rajagopal, E., Scherrer, Y., & Zampieri, M. (2021). Findings of the VarDial Evaluation Campaign 2021. *Proceedings of the Eighth Workshop on NLP for Similar Languages, Varieties and Dialects*, 1–11.

Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019, May). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding.

Federal Office of Topography swisstopo. (n.d.). The Swiss coordinates system.

Gaman, M., Hovy, D., Ionescu, R. T., Jauhiainen, H., Jauhiainen, T., Lindén, K., Ljubešić, N., Partanen, N., Purschke, C., Scherrer, Y., & Zampieri, M. (2020). A Report on the VarDial Evaluation Campaign 2020. *Proceedings of the 7th Workshop on NLP for Similar Languages, Varieties and Dialects*, 1–14.

Goldhahn, D., Eckart, T., & Quasthoff, U. (n.d.). Building Large Monolingual Dictionaries at the Leipzig Corpora Collection: From 100 to 200 Languages.

Hovy, D., & Purschke, C. (2018, October). Capturing Regional Variation with Distributed Place Representations and Geographic Retrofitting. In E. Riloff, D. Chiang, J. Hockenmaier & J. Tsujii (Eds.), *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing* (pp. 4383–4394). Association for Computational Linguistics.

Linder, L., Jungo, M., Hennebert, J., Musat, C., & Fischer, A. (2020, June). Automatic Creation of Text Corpora for Low-Resource Languages from the Internet: The Case of Swiss German.

Ljubešić, N., Samardžić, T., & Derungs, C. (2016). TweetGeo - A Tool for Collecting, Processing and Analysing Geo-encoded Linguistic Data. *Proceedings of COLING 2016, the 26th International Conference on Computational Linguistics: Technical Papers*, 3412–3421.

Lloyd, S. (1982). Least squares quantization in PCM. *IEEE Transactions on Information Theory*, *28*(2), 129–137.

Pfeiffer, J., Goyal, N., Lin, X., Li, X., Cross, J., Riedel, S., & Artetxe, M. (2022, July). Lifting the Curse of Multilinguality by Pre-training Modular Transformers. In M. Carpuat, M.-C. de Marneffe & I. V. Meza Ruiz (Eds.), *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies* (pp. 3479–3495). Association for Computational Linguistics.

Røyneland, U. (2009). Dialects in Norway: Catching up with the rest of Europe? *2009*(196-197), 7–30.

Scherrer, Y., & Ljubešić, N. (2020). HeLju@VarDial 2020: Social Media Variety Geolocation with BERT Models. *Proceedings of the 7th Workshop on NLP for Similar Languages, Varieties and Dialects*, 202–211.

Scherrer, Y., & Ljubešić, N. (2021). Social Media Variety Geolocation with geoBERT. *Proceedings of the Eighth Workshop on NLP for Similar Languages, Varieties and Dialects*, 135–140.

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is All you Need. *Advances in Neural Information Processing Systems*, *30*.